
fs | 23

A scalable, read-only, network filesystem with pervasive caching


John Salmon
D. E. Shaw Research
PDSW 2019 Work-In-Progress Talk
November 18, 2019

The Problem

- We have >15 petabytes of simulation data and >1 terabyte of code/binaries
 - Growing at ~10 terabytes/day and ~10 new software packages/versions/day.
 - **POSIX is non-negotiable for executables**
 - Read-only access is sufficient
 - Three geographically distributed data centers, remote workers, laptops,
...
 - NFS is a non-starter

The solution: fs123

- Read-only distributed POSIX filesystem
- How does it work?
 - Loosely-coupled client-server protocol built on HTTP
 - Client implements a Filesystem in USErspace (FUSE) filesystem
 - HTTP origin server exports a backend POSIX filesystem



- That's it!

The fs123 protocol: map low-level FUSE callbacks to HTTP

FUSE client gets callback from kernel:

```
fuse_lowlevel_ops::getattr(req, ino, fi)
```

FUSE client translates that into HTTP (or HTTPS):

```
GET http://server/ID/fs123/7/2/a/some/file
```

Origin server replies with:

```
HTTP/1.1 200
```

```
...
```

```
cache-control: max-age=86400, errno=0,  
uid=503, gid=503, mtime=1573923416, ...
```

FUSE client translates the reply into:

```
fuse_reply_attr(ino, &stat, timeout=86400)
```

The software

- A single client binary (no special permission required)

```
$ fs123p7 mount http://thesalmons.org:8888/ mtpt
```

- A single server binary (no special permission required)

```
$ fs123p7exportd –port 8888 –export-root=/public/stuff
```

- About 10k lines of C++, available on github (2-clause license):


<https://github.com/DEShawResearch/fs123>

- In production. Critical to our day-to-day scientific operations.

Why HTTP?

- Inherently wide-area
- Resilient on intermittent networks
- Standardized cache-management strategies
- Well understood by sysadmins

Well understood by sysadmins



Caching is essential for scalability

- Kernel caches
 - indispensable, require careful management
- Client-side disk caches
 - great for hiding network latency and coming back quickly after reboots
- Proxy caches (e.g., Varnish, Squid)
 - essential for wide-area operation

Caching would be easy if the data were immutable

- HTTP Cache-control (RFC7234) allows proxies to work read-only, mutable data
- fs123 adheres to RFC7234 for its kernel and disk cache
- RFC7234 is not quite enough:
 - Monotonic validator: “The file you’re asking about has changed since the last time you asked about it, so you (the client filesystem) should flush everything you have cached about its contents”.
 - Estale cookie: “The file you’re asking about (by name) has disappeared (inode) since the last time you asked about it, so any attempt to see more of it must fail with errno=ESTALE.”

Try it now

<https://github.com/DEShawResearch/fsI23>

IF you're comfortable running a static Linux binary from my personal URL:

```
$ wget https://thesalmons.org/fsI23/fsI23p7 && chmod +x fsI23p7  
$ mkdir mtpt c  
$ ./fsI23p7 mount -oFsI23CacheDir=c http://thesalmons.org:8888 mtpt  
# look around in mtpt: ls, find, cat, emacs (read-only)  
# if you feel lucky, and have devel versions of libevent, libcurl libsodium  
$ mkdir build; pushd build  
$ make -f .../mtpt/GNUmakefile  
# it's a fuse daemon. To shut it down, do:  
$ fusermount -u .../mtpt # or, in a pinch, pkill -9 fsI23p7
```